



Saving at epoch 55: ./cond_gpt/weights/more_heads_simplesplit_2layer_4head_16embd_32bs.pt

step_train_loss: 0.4115608334541321 train_step: 26000, learning_rate: 4e-05 | 94/471 [00:05<00:21, 17.54it/s]

epoch 56 iter 470: train loss 0.28608. lr 4.0000e-05: 100% | 471/471 [00:27<00:00, 16.91it/s]

test loss: %f 0.2539103700304931

epoch_valid_loss: 0.2539103700304931, epoch_train_loss: 0.39577958399070046, epoch: 56

step_train_loss: 0.4255453646183014 train_step: 26500, learning_rate: 4e-05 | 124/471 [00:07<00:20, 17.13it/s]

epoch 57 iter 470: train loss 0.41184. lr 4.0000e-05: 100% | 471/471 [00:28<00:00, 16.24it/s]

test loss: %f 0.25554204324506363

epoch_valid_loss: 0.25554204324506363, epoch_train_loss: 0.3967812232783899, epoch: 57

step_train_loss: 0.37357431650161743 train_step: 27000, learning_rate: 4e-05 | 152/471 [00:09<00:19, 16.38it/s]

epoch 58 iter 470: train loss 0.42348. lr 4.0000e-05: 100% | 471/471 [00:35<00:00, 13.22it/s]

test loss: %f 0.2562400299423146

epoch_valid_loss: 0.2562400299423146, epoch_train_loss: 0.3975178932814588, epoch: 58

step_train_loss: 0.37744730710983276 train_step: 27500, learning_rate: 4e-05 | 182/471 [00:11<00:17, 16.53it/s]

epoch 59 iter 470: train loss 0.40730. lr 4.0000e-05: 100% | 471/471 [00:29<00:00, 15.74it/s]

test loss: %f 0.2537996662112902

epoch_valid_loss: 0.2537996662112902, epoch_train_loss: 0.39384528520000966, epoch: 59

step_train_loss: 0.3576360046863556 train_step: 28000, learning_rate: 4e-05 | 210/471 [00:12<00:14, 17.62it/s]

epoch 60 iter 470: train loss 0.43753. lr 4.0000e-05: 100% | 471/471 [00:28<00:00, 16.68it/s]

test loss: %f 0.2489841934644951

epoch_valid_loss: 0.2489841934644951, epoch_train_loss: 0.39482868970579404, epoch: 60

Saving at epoch 60: ./cond_gpt/weights/more_heads_simplesplit_2layer_4head_16embd_32bs.pt

```
step_train_loss: 0.34282389283180237 train_step: 26000, learning_rate: 4e-05 | 94/471 [00:05<00:23, 16.26it/s]
epoch 56 iter 470: train loss 0.33361. lr 4.0000e-05: 100%| 471/471 [00:29<00:00, 15.97it/s]
test loss: %f 0.15094422830163307
epoch_valid_loss: 0.15094422830163307, epoch_train_loss: 0.2821874091136734, epoch: 56
step_train_loss: 0.22422336041927338 train_step: 26500, learning_rate: 4e-05 | 124/471 [00:07<00:20, 16.52it/s]
epoch 57 iter 470: train loss 0.25582. lr 4.0000e-05: 100%| 471/471 [00:29<00:00, 15.98it/s]
test loss: %f 0.14866365295536113
epoch_valid_loss: 0.14866365295536113, epoch_train_loss: 0.27892895201082696, epoch: 57
step_train_loss: 0.2607472240924835 train_step: 27000, learning_rate: 4e-05 | 152/471 [00:09<00:19, 16.53it/s]
epoch 58 iter 470: train loss 0.23728. lr 4.0000e-05: 100%| 471/471 [00:29<00:00, 16.01it/s]
test loss: %f 0.15319176369680548
epoch_valid_loss: 0.15319176369680548, epoch_train_loss: 0.2793911627776557, epoch: 58
step_train_loss: 0.30671072006225586 train_step: 27500, learning_rate: 4e-05 | 182/471 [00:10<00:17, 16.53it/s]
epoch 59 iter 470: train loss 0.27885. lr 4.0000e-05: 100%| 471/471 [00:29<00:00, 15.72it/s]
test loss: %f 0.14564041933923397
epoch_valid_loss: 0.14564041933923397, epoch_train_loss: 0.2805360517684062, epoch: 59
Saving at epoch 59: ./cond_gpt/weights/more_layers_simplesplit_4layer_2head_16embd_32bs.pt
step_train_loss: 0.26568254828453064 train_step: 28000, learning_rate: 4e-05 | 210/471 [00:13<00:16, 16.20it/s]
epoch 60 iter 470: train loss 0.29268. lr 4.0000e-05: 100%| 471/471 [00:29<00:00, 15.86it/s]
test loss: %f 0.14769485916169184
epoch_valid_loss: 0.14769485916169184, epoch_train_loss: 0.27897433601240695, epoch: 60
```

```
epoch_valid_loss: 0.2577916184123957, epoch_train_loss: 0.426970814380423, epoch: 56
Saving at epoch 56: ./cond_gpt/weights/my_experiment_simplesplit_2layer_2head_16embd_32bs.pt
step_train_loss: 0.36751508712768555 train_step: 26500, learning_rate: 4e-05 | 124/471 [00:04<00:12, 28.58it/s]
epoch 57 iter 470: train loss 0.44393. lr 4.0000e-05: 100% | 471/471 [00:17<00:00, 26.81it/s]
test loss: %f 0.2621100890748906
epoch_valid_loss: 0.2621100890748906, epoch_train_loss: 0.4276491145419467, epoch: 57
step_train_loss: 0.47926121950149536 train_step: 27000, learning_rate: 4e-05 | 151/471 [00:05<00:11, 28.09it/s]
epoch 58 iter 470: train loss 0.39307. lr 4.0000e-05: 100% | 471/471 [00:17<00:00, 26.63it/s]
test loss: %f 0.26155544255139695
epoch_valid_loss: 0.26155544255139695, epoch_train_loss: 0.4258403503464539, epoch: 58
step_train_loss: 0.39314109086990356 train_step: 27500, learning_rate: 4e-05 | 180/471 [00:06<00:10, 28.72it/s]
epoch 59 iter 470: train loss 0.36448. lr 4.0000e-05: 100% | 471/471 [00:17<00:00, 26.72it/s]
test loss: %f 0.26124161397511103
epoch_valid_loss: 0.26124161397511103, epoch_train_loss: 0.4227984719721881, epoch: 59
step_train_loss: 0.4268701374530792 train_step: 28000, learning_rate: 4e-05 | 211/471 [00:07<00:09, 28.20it/s]
epoch 60 iter 470: train loss 0.41615. lr 4.0000e-05: 100% | 471/471 [00:17<00:00, 26.86it/s]
test loss: %f 0.2590955774176796
epoch_valid_loss: 0.2590955774176796, epoch_train_loss: 0.42673711545148474, epoch: 60
```

step_train_loss: 0.35718634724617004 train_step: 23000, learning_rate: 4e-05 | 284/413 [00:10<00:04, 28.21it/s]
epoch 56 iter 412: train loss 0.45020. lr 4.0000e-05: 100% | 413/413 [00:16<00:00, 25.69it/s]
test loss: %f 0.2423201537002688
epoch_valid_loss: 0.2423201537002688, epoch_train_loss: 0.38743259730696966, epoch: 56
Saving at epoch 56: ./cond_gpt/weights/jump_split_addprim_jumpsplit_2layer_2head_16embd_32bs.pt
step_train_loss: 0.4419860541820526 train_step: 23500, learning_rate: 4e-05 | 371/413 [00:14<00:01, 27.83it/s]
epoch 57 iter 412: train loss 0.38042. lr 4.0000e-05: 100% | 413/413 [00:16<00:00, 24.98it/s]
test loss: %f 0.2449807777352955
epoch_valid_loss: 0.2449807777352955, epoch_train_loss: 0.3870527968712638, epoch: 57
epoch 58 iter 412: train loss 0.30742. lr 4.0000e-05: 100% | 413/413 [00:18<00:00, 22.04it/s]
test loss: %f 0.24279306083917618
epoch_valid_loss: 0.24279306083917618, epoch_train_loss: 0.38834499901489833, epoch: 58
step_train_loss: 0.3453693687915802 train_step: 24000, learning_rate: 4e-05 | 44/413 [00:01<00:15, 24.30it/s]
epoch 59 iter 412: train loss 0.34860. lr 4.0000e-05: 100% | 413/413 [00:16<00:00, 25.19it/s]
test loss: %f 0.24373750129471655
epoch_valid_loss: 0.24373750129471655, epoch_train_loss: 0.3863854018236188, epoch: 59
step_train_loss: 0.3549240529537201 train_step: 24500, learning_rate: 4e-05 | 133/413 [00:05<00:11, 25.02it/s]
epoch 60 iter 412: train loss 0.36386. lr 4.0000e-05: 100% | 413/413 [00:17<00:00, 24.15it/s]
test loss: %f 0.24116106979224994
epoch_valid_loss: 0.24116106979224994, epoch_train_loss: 0.3859563932436142, epoch: 60
Saving at epoch 60: ./cond_gpt/weights/jump_split_addprim_jumpsplit_2layer_2head_16embd_32bs.pt

```
step_train_loss: 0.1917940378189087 train_step: 26000, learning_rate: 4e-05 | 94/471 [00:04<00:16, 23.16it/s] ✓ 43
epoch 56 iter 470: train loss 0.24639. lr 4.0000e-05: 100% | 471/471 [00:21<00:00, 21.64it/s]
test loss: %f 0.07889651907783635
epoch_valid_loss: 0.07889651907783635, epoch_train_loss: 0.19815958814889256, epoch: 56
step_train_loss: 0.21368414163589478 train_step: 26500, learning_rate: 4e-05 | 122/471 [00:05<00:15, 23.05it/s]
epoch 57 iter 470: train loss 0.14465. lr 4.0000e-05: 100% | 471/471 [00:21<00:00, 21.57it/s]
test loss: %f 0.07640211487997253
epoch_valid_loss: 0.07640211487997253, epoch_train_loss: 0.19602822563007885, epoch: 57
step_train_loss: 0.22735927999019623 train_step: 27000, learning_rate: 4e-05 | 153/471 [00:06<00:13, 23.04it/s]
epoch 58 iter 470: train loss 0.19718. lr 4.0000e-05: 100% | 471/471 [00:21<00:00, 21.70it/s]
test loss: %f 0.07591710935504932
epoch_valid_loss: 0.07591710935504932, epoch_train_loss: 0.19713622026319463, epoch: 58
step_train_loss: 0.2261740118265152 train_step: 27500, learning_rate: 4e-05 | 182/471 [00:08<00:12, 22.62it/s]
epoch 59 iter 470: train loss 0.17536. lr 4.0000e-05: 100% | 471/471 [00:21<00:00, 21.69it/s]
test loss: %f 0.07428785294013203
epoch_valid_loss: 0.07428785294013203, epoch_train_loss: 0.1961450825134913, epoch: 59
Saving at epoch 59: ./cond_gpt/weights/larger_embd_simplesplit_2layer_2head_32embd_32bs.pt
step_train_loss: 0.21199221909046173 train_step: 28000, learning_rate: 4e-05 | 210/471 [00:09<00:11, 22.62it/s]
epoch 60 iter 470: train loss 0.14677. lr 4.0000e-05: 100% | 471/471 [00:21<00:00, 21.55it/s]
test loss: %f 0.07441298642529631
epoch_valid_loss: 0.07441298642529631, epoch_train_loss: 0.19507814562561152, epoch: 60
```